

SC4031 IoT Course Project

JESSICA DANIELLA GIRSANG
(U2221579J)



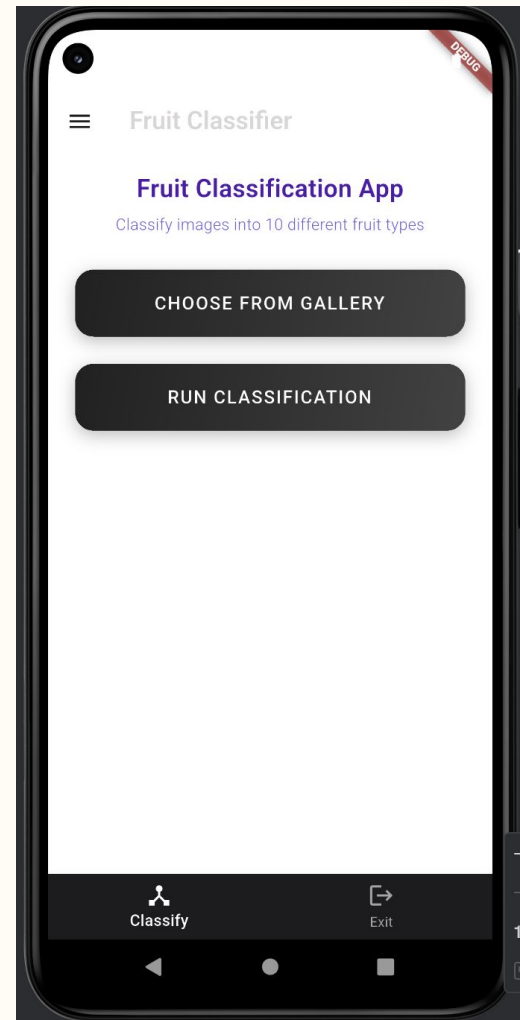
Fruit Classification App

1 Classify fruit (**10 classes**):

Apple	Avocado
Banana	Mango
Pineapple	Strawberries
Kiwi	Cherry
Orange	Watermelon

2 Image Input : User Gallery

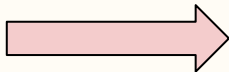
3 Functions: Quick fruit recognition (personal, dietary, etc.)



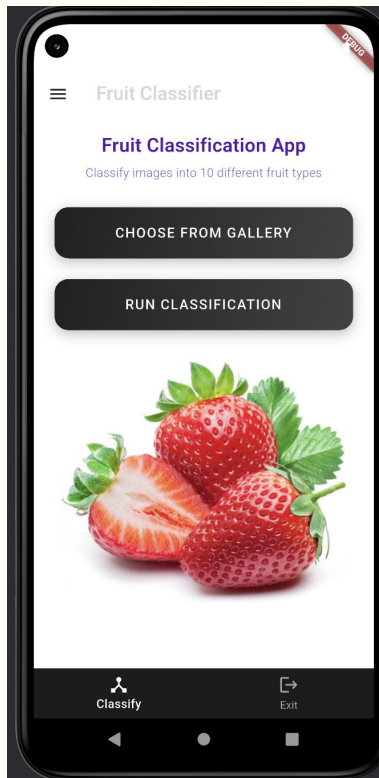
Architecture



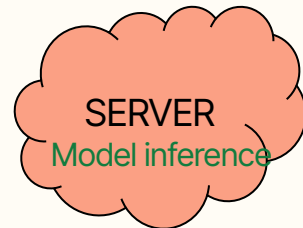
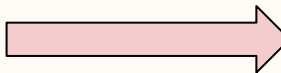
Upload Image



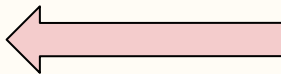
DISPLAY RESULTS



HTTP POST request
(Send Image)



HTTP returns
prediction results



Model

DATASET: <https://www.kaggle.com/datasets/karimabdulnabi/fruit-classification10-class>

```
[*]: model.fit(train_ds, epochs=15, callbacks=callbacks, validation_data=val_ds)
```

Epoch 1/15

58/58 ————— 0s 16s/step - accuracy: 0.2589 - loss: 1.9617

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.

58/58 ————— 969s 17s/step - accuracy: 0.2604 - loss: 1.9585 - val_accuracy: 0.0897 - val_loss: 2.3058

Epoch 2/15

58/58 ————— 0s 16s/step - accuracy: 0.4466 - loss: 1.5672

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.

58/58 ————— 959s 16s/step - accuracy: 0.4471 - loss: 1.5655 - val_accuracy: 0.1050 - val_loss: 2.3255

Epoch 3/15

58/58 ————— 0s 16s/step - accuracy: 0.4723 - loss: 1.4819

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.

58/58 ————— 969s 17s/step - accuracy: 0.4723 - loss: 1.4819 - val_accuracy: 0.1201 - val_loss: 2.3753

Epoch: 15+5


Optimizer: Adam

BatchSize = 32



APP DEMO

Local Inference

Marking Criteria (1)

- Collect user input (10)
 - **If** load input from storage (5) 
 - **ELSE** collect real-time input by touch screen, microphone, camera, built-in sensors (10)

Marking Criteria (2)

- Infer locally and display result (15)
 - **If** your app offloads the execution to cloud, (10) 
 - **ELSE** run heuristic algorithm (not neural network) on the mobile app (5)
- Display the inference result in real time by screen or synthetic voice (5) 

Marking Criteria (3)

- Run on emulated/physical IoT device (15)
 - **If** the program runs natively on a desktop or laptop computer (5)
 - **ELSE** the program runs on an emulated or a physical IoT device (15)
 - iOS emulator provided by Xcode
 - Android emulator
 - Real smartphone
 - Raspberry Pi + add-on sensors (camera, microphone, etc)



Cloud Inference

Marking Criteria (4)

- Run inference in cloud virtual machine (10)
 - **If** deploy server program on a cloud virtual machine, e.g., Azure. (10)
 - **ELSE** deploy server program on your own computer (10)



Marking Criteria (5)

- Communication btw IoT device and cloud (20)
 - Send the user input from the mobile app to the cloud for inference (10)
 - Send the inference result from the cloud to the mobile app (10)



Advanced Task

Marking Criteria (6)

- Train your own model (10)
 - **If** train the used ML model by yourself (10)
 - Training program should be in the code package
 - Training results (e.g., accuracy) in a readme file
 - The training data can be any publicly available dataset
 - **ELIF** use downloaded pre-trained model (7)
 - **ELSE** use heuristic algorithm (not neural network) (3)



Marking Criteria (7)

- Support multiple concurrent users (10)
 - Demonstrate multiple IoT devices can use the cloud service simultaneously



A blue rectangular button with a folded corner effect on the top right, revealing an orange layer underneath. The text "Thank You" is centered on the blue surface.

Thank You